APOLLO
RESEARCH

Research Scientist/Engineer (Evaluations)

APPLY FOR THIS JOB

London & San Francisco
Scheming Research / Full-Time / On-Site

Application deadline: We are conducting interviews actively and aim to fill this role as soon as we find someone suitable.

ABOUT THE OPPORTUNITY

We're looking for **Research Scientists/Engineers** for our pre-deployment team to work on Training-Run Assessments (TRAs). You will design and build automated pipelines for assessing whether egregious misalignment or scheming are emerging at any point of frontier post-training.

This will involve evaluating and red-teaming of checkpoints at various stages of post-training as well as automated analysis of post-training data.

You will get to work with frontier labs like OpenAI, Anthropic, and Google DeepMind and be among the first to interact with new models before anyone else. Our ideal candidate loves rigorously testing frontier AI models, and enjoys building efficient pipelines for automated analysis.

KEY RESPONSIBILITIES

- **Run and own pre-deployment engagements:** We run a pre-deployment evaluation campaign with a frontier AI lab approximately every two weeks with thousands of runs across hundreds of distinct environments. We explore behaviors learned during training and check for undesirable behaviors like alignment faking, perform targeted follow-up experiments/red-teaming, and report our findings to the frontier AI lab we're working with.
- **Develop methodology for training-run assessments:** in-between campaigns we improve our methodology, which might mean implementing new evals or building infrastructure for automated red-teaming.

KEY REQUIREMENTS

- We don't require a formal background or industry experience and welcome self-taught candidates.
- **Software engineering skills:** Our entire stack uses Python. We're looking for candidates with strong software engineering experience. Ideally, you have experience shipping and maintaining production Python code, and know how to factor messy problems into clean abstractions that others can use and extend.
- **Data Analysis & Pattern Recognition:** You can extract signal from large, messy datasets. You're comfortable with quantitative analysis and know when qualitative assessment is more appropriate. You can identify anomalies and unexpected model behaviors.
- **Writing and communication:** You succinctly convey qualitative and quantitative findings to a technical and non-technical audience.
- **AI power-user:** You're capable of using AI to accelerate your work, technical or otherwise. You have experience using different models, know which ones to use for which tasks, when not to use AI, and always experiment with new AI workflows.

NICE TO HAVE

- Experience thinking about AI risk topics like scheming and metagaming.
- Knowledge of LLM post-training: topics like RLHF, reasoning training, supervised fine-tuning, on-policy distillation, etc.
- We are using Inspect as our primary evals framework, and we value experience creating evals with it or similar frameworks like Harbor.

We want to emphasize that people who feel they don't fulfill all of these characteristics but think they would be a good fit for the position, nonetheless, are strongly encouraged to apply. We believe that excellent candidates can come from a variety of backgrounds and are excited to give you opportunities to shine.

BENEFITS

- **This role offers market competitive salary, equity, and competitive benefits.**
- Salary: 100k - 200k GBP (~150k - 270k USD). We will be looking to meaningfully raise salaries soon.
- Flexible work hours and schedule
- Unlimited vacation
- Unlimited sick leave
- Up to 6 months of paid parental leave
- Comprehensive health, dental and vision insurance
- Retirement savings with competitive employer matching (e.g. 401(k) for US employees)
- Lunch, dinner, and snacks are provided for all employees on workdays
- Paid work trips, including staff retreats, business trips, and relevant conferences
- A yearly $1,000 (USD) professional development budget
- Relocation support and visa fees (if applicable)

LOGISTICS

- **Time Allocation**: Full-time
- **Location**: This is an in-person role working out of our London or San Francisco office. We offer flexible working hours and wfh arrangements.
- **Visa sponsorship:** We sponsor visas in both the UK and US. Sponsorship isn't guaranteed for every role or candidate, but if we make you an offer, we'll work with you to find the right visa route

ABOUT APOLLO RESEARCH

The rapid rise in AI capabilities offer tremendous opportunities, but also present significant risks.At Apollo Research, we're primarily concerned with risks from Loss of Control, i.e. risks coming from the model itself rather than e.g. humans misusing the AI. We're particularly concerned with deceptive alignment / scheming, a phenomenon where a model appears to be aligned but is, in fact, misaligned and capable of evading human oversight.

We work on the detection of scheming (e.g., building evaluations and novel evaluation techniques), the science of scheming (e.g., model organisms and the study of scaling trends), and scheming mitigations (e.g., control). We closely work with multiple frontier AI companies, e.g. to test their models before deployment and collaborate on fundamental research.

At Apollo, we aim for a culture that emphasizes truth-seeking, being goal-oriented, giving and receiving constructive feedback, and being friendly and helpful. If you're interested in more details about what it's like working at Apollo, you can find more information here.

ABOUT THE TEAM

The current evals team consists of Jérémy Scheurer,Alex Meinke,Bronson Schoen, Felix Hofstätter,Axel Højmark,Teun van der Weij,Alex Lloyd, Mia Hopman, Dylan Bowman and Ezra Newman . Alex Meinke coordinates the research agenda with guidance from Marius Hobbhahn, though team members lead individual projects. You will mostly work with the evals team as well as our team of software engineers, but you will likely sometimes interact with the governance team to translate technical knowledge into concrete recommendations. You can find our full team here.

Equality Statement: Apollo Research is an Equal Opportunity Employer. We value diversity and are committed to providing equal opportunities to all, regardless of age, disability, gender reassignment, marriage and civil partnership, pregnancy and maternity, race, religion or belief, sex, or sexual orientation.

How to apply: Please complete the application form with your CV. The provision of a cover letter is optional but not necessary. Please also feel free to share links to relevant work samples.

About the interview process: Our multi-stage process includes a screening interview, a take-home test (approx. 2.5 hours), 3 technical interviews, and a final interview with Marius (CEO). The technical interviews will be closely related to tasks the candidate would do on the job. There are no LeetCode-style general coding interviews. If you want to prepare for the interviews, we suggest working on hands-on LLM evals projects (e.g. as suggested in our starter guide), such as building LM agent evaluations in Inspect.

Your Privacy and Fairness in Our Recruitment Process: We are committed to protecting your data, ensuring fairness, and adhering to workplace fairness principles in our recruitment process. To enhance hiring efficiency, we use AI-powered tools to assist with tasks such as resume screening. These tools are designed and deployed in compliance with internationally recognized AI governance frameworks.Your personal data is handled securely and transparently. We adopt a human-centred approach: all resumes are screened by a human and final hiring decisions are made by our team. If you have questions about how your data is processed or wish to report concerns about fairness, please contact us at info@apolloresearch.ai.

APPLY FOR THIS JOB

Jobs powered by LEVER