Program Officer, AI Model Safety

AI Resilience · San Francisco

Apply now ↗

About OpenAI Foundation

AI is already changing how people work, learn, and access care. It has the potential to unlock extraordinary benefits—faster medical breakthroughs, accelerated scientific discovery, more personalized services in healthcare and education, new tools for creativity and invention, higher productivity and economic growth, improved public services like transportation systems, and so much more.

But building powerful systems to benefit humanity is only a part of the work of our mission. Advanced AI models will also present new challenges that are already surfacing, and we need to be prepared to identify these challenges and develop solutions to address them.

These are the two dimensions of what we're building the Foundation to do. We aim to enable the use of AI to find solutions to humanity's hardest problems, transform what people are capable of, and deliver real benefits in people's lives—while working hard with partners to be ready for new challenges, and to help make society resilient, as AI advances.

This work is just beginning. Over the next year, as we quickly ramp up, the Foundation expects to invest at least $1 billion across life sciences and curing diseases, jobs and economic impact, AI resilience, and community programs. This includes early investments toward our previously announced $25 billion commitment to curing diseases and AI resilience.

About the team

AI Resilience helps society absorb, adapt to, and benefit from increasingly capable AI systems, with work across biosecurity, model safety, cybersecurity, children and youth, and the future of the economy. The team is small and fast-moving with many workstreams, candidates, advisors, grantees, board touchpoints, and special projects in flight.

Within AI Resilience, the AI Model Safety program works to make AI systems safer by default. By AI model safety we mean the technical work of understanding and improving how models behave (testing and evaluating them independently, red-teaming them, researching how to interpret, align, and control them) so that increasingly capable systems can be trusted with increasingly consequential tasks. The program supports this through funding independent testing and evaluations, new and stronger industry standards, and foundational research that helps avoid safety issues or detect and address them early.

About the role

We are looking for a founding grantmaker for the Foundation's AI Model Safety program: someone with deep technical fluency in how frontier models are evaluated, audited, and made safer, who can turn that into a grant portfolio and a stronger field.

You will help with the program's grantmaking end to end (investigating opportunities across independent evaluations, standards, and foundational safety research) and do the proactive work, such as identifying the organizations that should exist, finding the people to build them, and helping the field scale as fast as the models do. You will develop the program's strategy together with the rest of AI Resilience, and work closely with the AI Resources team, whose AI resources your grantees will increasingly leverage.

This is an early, zero-to-one role with significant independence. You would report to Lennart Heim, Director of AI Resources, for the moment. We are open to mid-career and senior experience levels.

In this role, you will

- Run the program's grantmaking across its domains (independent testing and evaluations, industry standards, foundational safety research, AI control, and other new ideas) from sourcing through investigation to getting resources into grantees' hands.
- Run grant investigations and make grant recommendations to the board.
- Proactively identify the most important projects and organizations that need to exist, and make them happen: scope priority projects, find or develop the right founders, and actively build new initiatives, seeding them if required.
- Work with the AI Resources team to route AI resources to safety grantees and help them absorb it productively.
- Develop the program's strategy with the team lead: what we should fund, what others in the ecosystem are doing, and where the gaps are.
- Build and maintain relationships across the safety ecosystem: independent safety organizations, AI safety institutes, academic groups, other funders (philanthropic and VC alike), and AI companies' safety teams.

You might thrive in this role if you

- Have at least 3 years of experience in AI safety, ideally including hands-on research in model safety or evaluations, whether in a research role, on a safety or evals team, in grantmaking, or helping start an organization.
- Have deep, current fluency in AI model safety: evaluations, red-teaming, auditing, interpretability, control, and alignment research.
- Follow frontier AI capabilities and risks closely enough to have independent judgment about them, with views on the risks that matter most and on what each part of the ecosystem should do.
- Have a view on where AI is going, which risks it brings, and what to do about them: your own takes on how capabilities will develop, and what that means for the safety work that will be needed.
- Excel at owning ambitious, loosely scoped projects: you find the right people, convince them to start new work, and support their success.
- Have a relentless prioritization mindset: you evaluate opportunities on impact, tractability, and neglectedness, and act on the ranking.
- Write and communicate: compelling, transparent documents that explain your reasoning and appropriately highlight risks, uncertainties, and downsides internally.
- Are deeply mission-aligned and excited to help ensure AGI benefits all of humanity.

It would be great if you also have

- Grantmaking, VC, or philanthropic experience, especially seeding organizations or sub-fields.
- A research track record in evaluations, interpretability, control, or adjacent technical safety areas.
- Existing relationships across safety organizations, AI safety institutes, and AI company safety teams.

Compensation

The annual base salary range for this role is $220,000 to $300,000 USD.

This range reflects our good-faith estimate of the base salary we reasonably expect to pay for this position upon hire. Actual base salary will be determined based on job-related factors, including the candidate's relevant experience, skills, expertise, the scope of the role, and work location, as applicable.

This salary range does not include benefits or any other compensation for which the role may be eligible.

Equal Opportunity

We are an equal opportunity employer and do not discriminate on the basis of race, religion, national origin, gender, sexual orientation, age, veteran status, disability or any other legally protected status. Pursuant to the San Francisco Fair Chance Ordinance, we will consider qualified applicants with arrest and conviction records.

Apply now ↗

© 2015–2026

Terms of Use Privacy Policy Cookie Policy English United States