Economics

A research agenda for the Economic Futures Research Fund

Jul 22, 2026

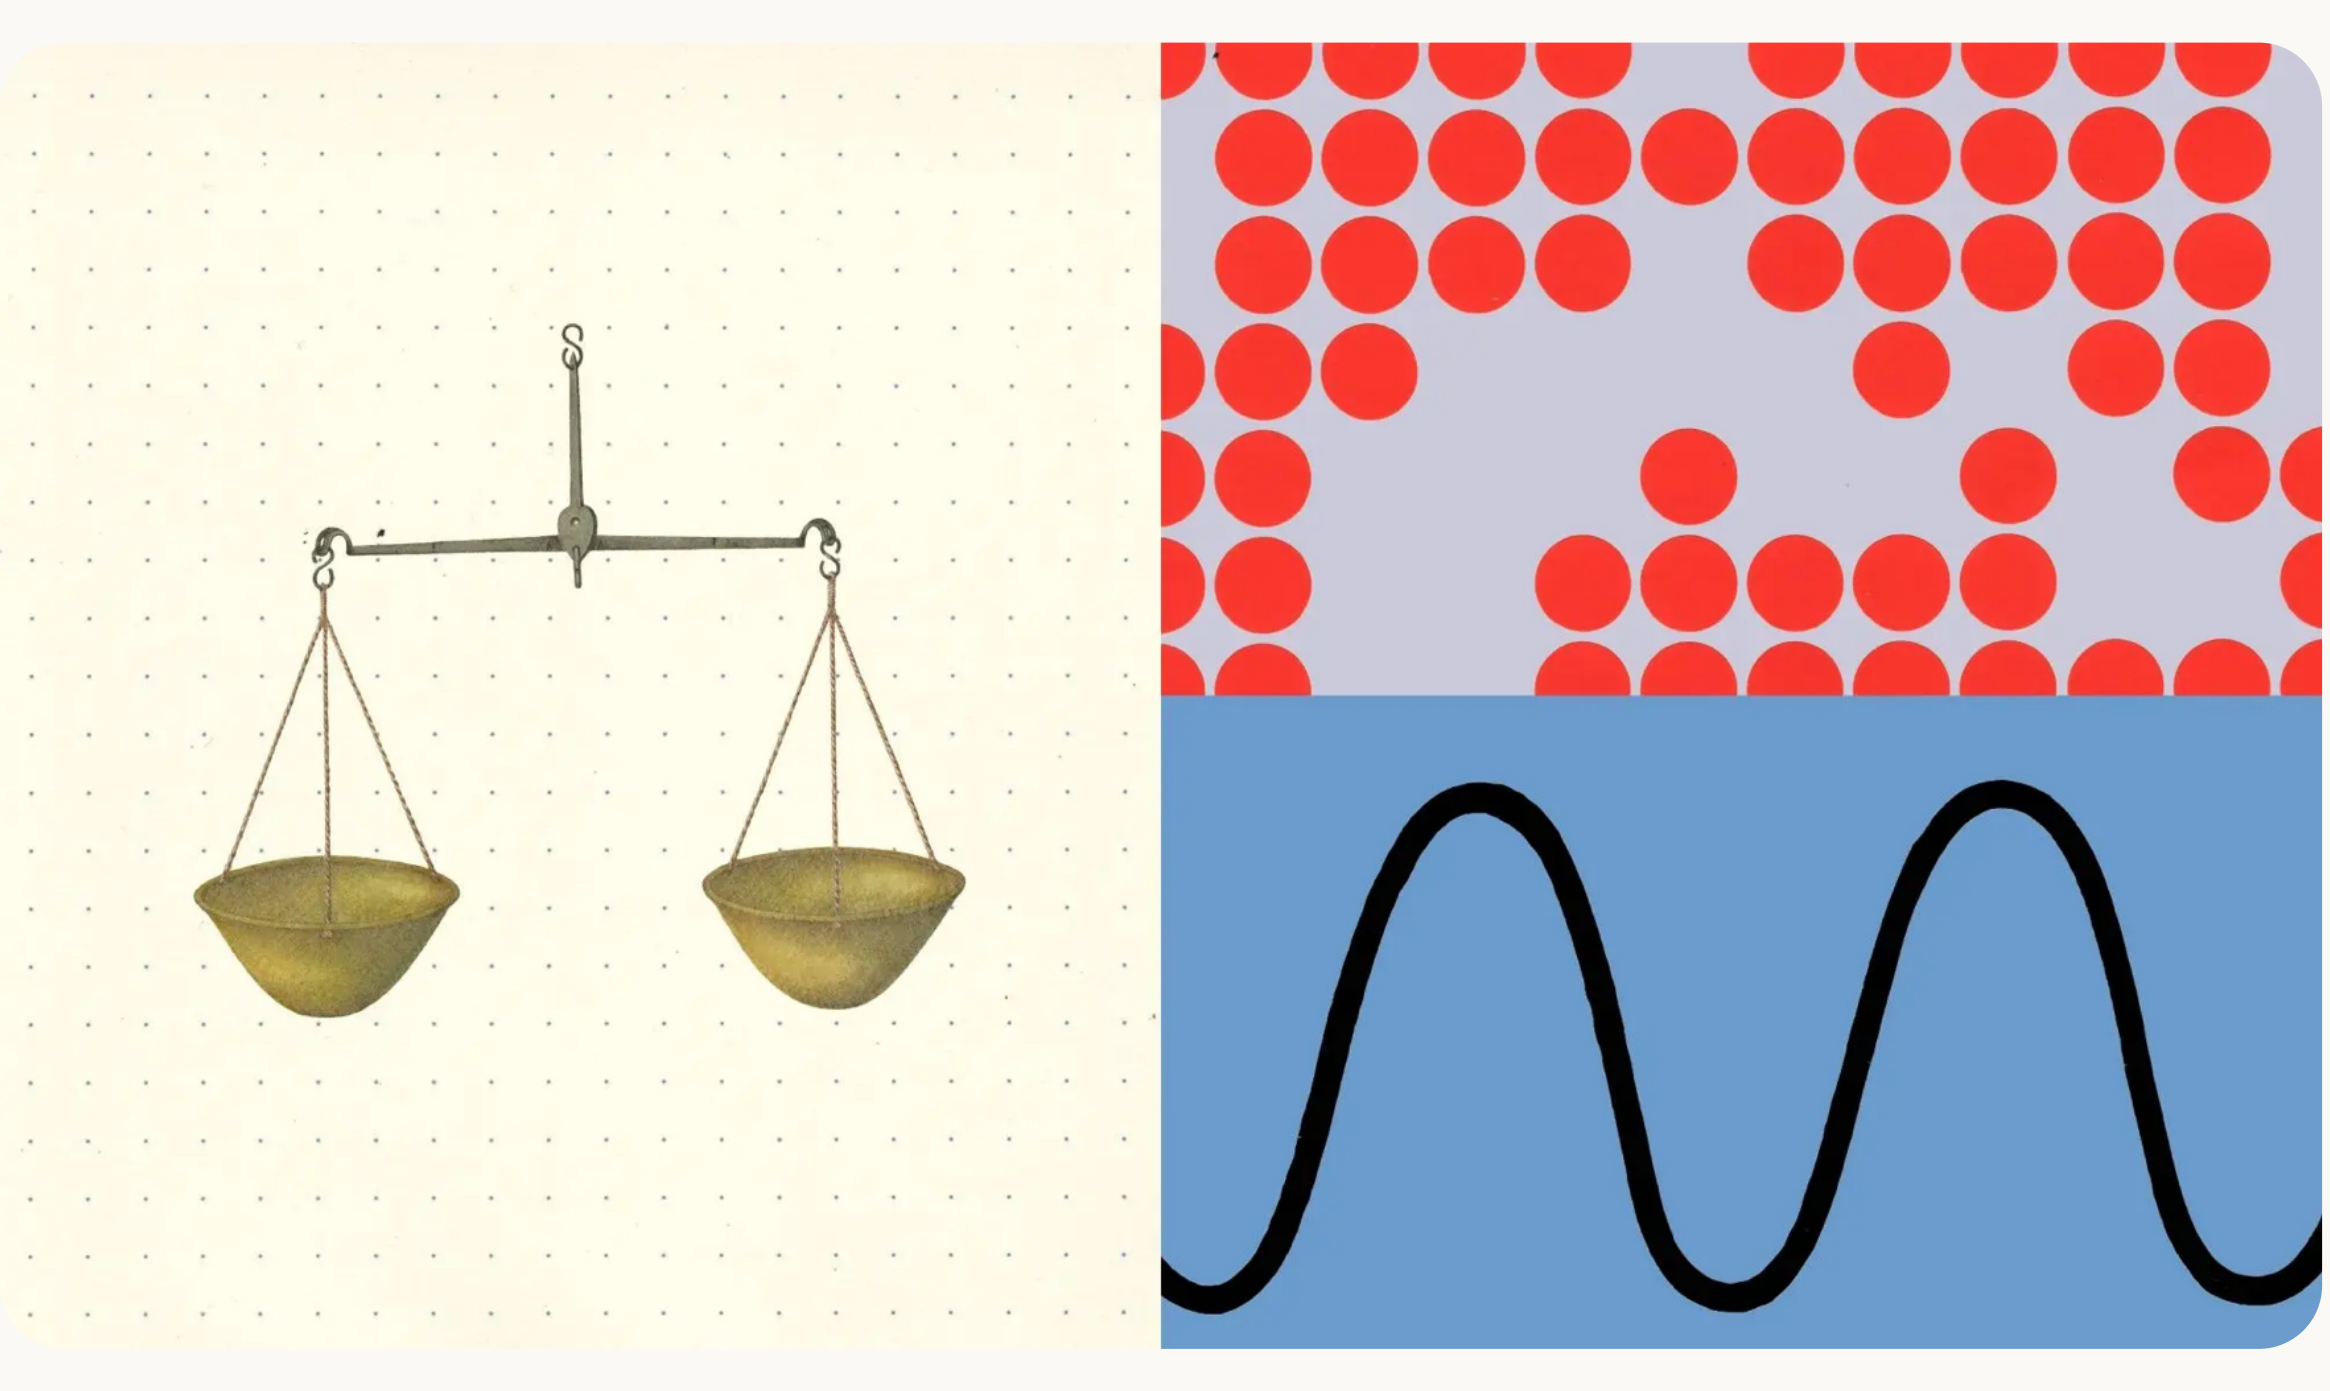

We're sharing the research agenda for the Anthropic Economic Futures Research Fund. We're committing $200 million to the fund to support ambitious external research on interventions to prepare society for the economic impacts of AI.

With the research the Fund supports, we want to study what programs could make the economy more flexible and resilient, ensure the benefits of AI are shared, and minimize the harm that AI-driven disruption could cause.

In the fund, we'll prioritize five research areas:

1. Shaping AI's impact on workers at the firm and workplace level
2. Equipping people to navigate AI-driven transitions
3. Modernizing income support for AI-driven displacement
4. Building worker stakes in AI-driven growth before disruption arrives
5. Generating new evidence on public investments

What is the fund?

AI capabilities continue to improve. But we don't yet know how quickly AI will diffuse throughout the economy while also becoming more capable, and what the economic effects will be. In our Economic Policy Framework (EPF), published in June, we proposed programs and policies for a range of scenarios. But we need more empirical evidence on which interventions might actually work in an AI-transformed economy—which ones make the economy more flexible and resilient, and spread the gains broadly. In the face of this uncertainty, our aim is to build this evidence base so that workers, firms, and governments have room to adapt. This $200 million fund will support external research on interventions proposed in the EPF and on other open questions. We may be entering a moment without historical precedent, where the most promising solutions are ones nobody has tried yet. We're willing to fund creative and ambitious pilots that can provide guidance on questions where randomized control trials alone might only provide incremental evidence.

This is a significant evolution of our Economic Futures program, launched a year ago. We're updating our focus to ambitious projects and large grants, because we think it's where we can have the highest impact. We've always thought that it's important to fund big external research bets; this shift will let us do so. We also learned through the Economic Futures program that it's hard for us to scale capacity to manage many small grants at once. In addition to large-scale RCTs and pilots, we're also interested in working with partners that could scale up a program of effective small-scale pilots.

The kinds of research we're funding

Fundamentally, we want to fund the most ambitious proposals possible. We aim to fund large-scale RCTs or ambitious, creative pilots or program evaluations that expand our shared understanding of what shows promise and in which contexts, fill gaps where evidence is thin, and inspire new solutions. The fundable directions below lay out a set of possibilities, but we know that we have not come up with all the good ideas in this space. We welcome proposals that may not be captured below.

AI could transform society faster than traditional research funding and publication cycles can keep pace with. We're looking to partner with research organizations that are willing to share what they're learning publicly at key milestones because a signal that arrives early enough to act on can be worth more than an answer that arrives too late. We're especially interested in pilots that can be scaled up dramatically if they show promise.

This is a global fund. The funding directions that we've outlined below are somewhat US-centric, in part because we're headquartered in San Francisco, and Claude is used more in the US than any other country. But the need to prepare for disruption will be necessary worldwide, and we expect to fund projects in a way that reflects that.

We plan to primarily fund projects in the $5-30 million range, though we're flexible upward for well-scoped, high-potential-impact projects. Based on what we learned from the Economic Futures program, and the ambition and scale we're seeking in proposals, we won't directly fund anything below $1 million from this fund.

We'll accept proposals from accredited universities and other degree-granting institutions, from independent research institutes and policy research organizations, and from nonprofits with a track record of running field experiments at scale. Individual researchers may serve as principal investigators on proposals made by their institutions, but we won't consider proposals from individuals applying in their personal capacity.

We're more likely to fund projects that fit one of our research priorities, but we welcome ambitious proposals outside them, as long as they're calibrated to the scale of the problem and opportunity. See our request for proposals and apply here.

Our five research priorities

1. Shaping AI's impact on workers at the firm and workplace level

AI's impact on the labor market depends on the systems, workplaces, training protocols, and institutional choices that are built around it. The existing evidence on AI's integration in the workplace is observational and short-term. Field experiments can help us understand which collaborative patterns develop human expertise alongside AI, how organizational design choices affect both productivity and who captures the gains, and what difference worker voice makes in those design choices.

Without this evidence, both firm-level decisions and policy levers like incentives for worker augmentation, retention tax credits, employer co-investment requirements, or apprenticeship programs will be poorly informed.

Fundable directions include:

- Field experiments randomizing AI systems and AI integration designs at the firm or team level, including comparisons of designs co-developed with workers and worker organizations against top-down approaches.
- Estimates of the impact of organizational choices around AI workplace integration and usage on the incidence of AI productivity gains.
- Evaluations of retention tax credits and employer co-investment requirements.

2. Equipping people to navigate AI-driven transitions

The evidence on retraining and job placement is mixed, and it may not generalize to AI-induced economic disruption and rapid structural transformation.

Fundable directions include:

- Evaluations of innovative skill retraining, job placement, licensing reform, and sectoral transition packages, including newer AI-enabled matching, credentialing, and learning models, and the bundling of income support with intensive reemployment services, retraining, and relocation assistance.
- Field experiments on the early-career and professional pipeline, e.g., what apprenticeship, mentorship, or rotational models can build expertise if junior tasks are absorbed by AI.
- Evaluations of curriculum and educational delivery models in K-12 and higher education that aim to prepare students for a transformed labor market, including longitudinal pilots that link educational interventions to later labor market outcomes.
- Tests of ambitious mobility instruments, for example paid leave tied to retraining programs and portable benefits that follow workers across employers.

There's existing evidence on many such efforts, including some especially effective sectoral training programs. We want to find out whether promising programs could scale quickly across a broader population. For example, a large-scale "fire drill" where selected programs are scaled up rapidly for job seekers in a given state could provide evidence on how well these programs work in the face of major disruption.

3. Modernizing income support for AI-driven displacement

Like similar insurance programs around the world, the US system for supporting displaced workers is built almost entirely around the assumption that joblessness is temporary. AI may lead to displacement that is broader and more persistent. In that scenario, we'll need instruments calibrated to a new equilibrium, one with no modern precedent.

Fundable directions include:

- Unemployment Insurance (UI) reforms suited to AI-driven displacement, including alternative eligibility thresholds, automatic extension triggers linked to industry or occupation, and integration of UI with wage insurance, retraining, or other transition supports.
- Basic needs relief for workers who exhaust UI, never qualified, or are persistently underemployed.
- Longer-duration unconditional income pilots at livable levels, designed to speak to scenarios where income and work are decoupled for sustained periods of time, with analyzed outcomes spanning not only labor supply and consumption but also wellbeing, family stability, child development, civic participation, and how recipients structure their time.

4. Building worker stakes in AI-driven growth before disruption arrives

In unprecedented scenarios where AI delivers large aggregate gains, those gains may not be broadly shared by default. In the EPF, we discuss universal pre-distributive capital accounts and adjacent mechanisms, like equity-sharing, AI-sector dividends, and public ownership stakes. But these mechanisms have limited direct empirical precedent at scale, and they also need a funding source. Many proposals to generate revenue exist, including taxing AI-driven returns through corporate, capital gains, or token taxes. But we lack evidence on who would bear the economic incidence of such taxes, and how different designs would affect collected revenue and adoption.

Fundable directions include:

- RCTs testing the design of pre-distributive capital accounts at scale.
- Pilots testing equity-sharing or dividend-style mechanisms, including community-level pilots where AI infrastructure or AI-using firms generate direct, ongoing returns to local residents.
- Evaluations comparing different mechanisms for raising and distributing revenue—which tax base (corporate profits, capital gains, compute, automation taxes, etc.) and which mechanism (pre-distributive accounts, equity stakes, dividends, or equivalent direct transfers) lead to the best labor market and household outcomes.

5. Generating new evidence on public investments

The EPF calls for both modernizing the income safety net and substantially expanding public investment in human- and community-facing work. Policymakers need a consistent way to compare these instruments against one another, and against direct transfers. This research would generate evidence on what forms of spending generate the most public benefit, especially in sectors that might be undervalued by the private market.

Fundable directions include:

- Large-scale pilots that directly fund human- and community-facing service positions (in e.g., teaching, after-school programming, libraries, community health, parks, infrastructure, the arts), measuring outcomes including employment levels, educational attainment, crime, and wellbeing.
- Pilots broadening access to AI-enabled public services (legal aid, medical guidance, financial advice) for underserved populations, testing whether such investments can narrow the divide in access.
- Guaranteed-jobs pilots for displaced or long-term unemployed workers, in which participants are offered employment in public good roles in the spirit of the Civilian Conservation Corps but spanning a broader range of roles.
- Place-based interventions in communities most exposed to AI-driven displacement or hosting major AI infrastructure build-outs, including bundled investments in workforce, public services, infrastructure, and amenities, and pilots of regional development authorities that coordinate these investments under unified governance.

Learn more about the RFP and apply here.

Related content

Developing Enterprise Frontier Safeguards with our customers

Read more →

Improving our alignment and security efforts

On July 30, we reported three incidents in which Claude models gained unauthorized access to real computer systems. We are conducting an in-depth analysis of both incidents, and planning to work with METR for an independent review. In the meantime, we're sharing some of the changes we've made over the past month.

Read more →

Previewing the Model Hardware Standard

We're opening a research preview of the Model Hardware Standard (MHS), a shared specification for AI agents to safely operate physical devices, to a first group of scientific research labs and advanced manufacturers.

Read more →

Products
Claude
Claude Code
Claude Code Enterprise
Claude Cowork
@Claude
Claude Design
Claude Science
Claude Security
Claude in Chrome
Claude for Microsoft 365
Skills
Download app
Pricing
Log in to Claude

Models
Mythos
Fable
Opus
Sonnet
Haiku

Solutions
AI agents
Code modernization
Coding
Customer support
Cybersecurity
Enterprise
Financial services
Government
Healthcare
Higher education
K-12 teachers
Legal
Life sciences
Nonprofits
Small business

Claude Platform
Overview
Developer docs
Pricing
Ecosystem
Marketplace
Regional compliance
Claude on AWS
Google Cloud
Microsoft Foundry
Console login

Resources
Blog
Claude partner network
Community
Connectors
Courses
Customer stories
Engineering at Anthropic
Events
Plugins
Powered by Claude
Service partners
Tutorials
Use cases

Programs
Startups
Research Labs

Help and security
Availability
Status
Support center

Company
Anthropic
Careers
Leadership
Policy
Economic Futures
Research
News
Claude's Constitution
Claude Corps
Keep thinking
Policy on the AI Exponential
Responsible Scaling Policy
Security and compliance
Transparency

Terms and policies
Privacy choices
Privacy policy
Consumer health data privacy policy
Responsible disclosure policy
Terms of service: Commercial
Terms of service: Consumer
Terms of Service: US K-12
Data Processing Agreement: US K-12
Usage policy

© 2026 Anthropic PBC. Services in the EU are provided by Anthropic Ireland Limited.