Benchmarks

A registry of 85 AI benchmarks covering mathematics, software engineering, agentic workflows, games, and more.

Filter

Creator

☐ Made by Epoch (9)

Review verdict

How we review

☐ Verified (4)
☐ Flawed (9)
☐ Not enough info (2)

Highest score

☐ No percentage score (12)
☐ Under 30% (4)
☐ 30–70% (19)
☐ Over 70% (45)

Administered by

☐ Epoch AI (14)
☐ Benchmark creator (52)
☐ Model developer (18)

Domain

☐ Agent (29)
☐ Software engineering (22)
☐ Long context (10)
☐ Multimodal (10)
☐ Science (9)
☐ Mathematics (8)
☐ Language understanding (8)
☐ Games (7)
+ More

Search

Sort

85 results

Epoch Capabilities Index (ECI)

By Epoch

An index aggregating many different benchmarks into a single, general capability scale.

Epoch administered

View benchmark

FrontierMath: Open Problems

By Epoch

A collection of unsolved problems from research mathematics, curated to be especially interesting and difficult.

Mathematics | Epoch administered

View benchmark

FrontierMath Erdős

By Epoch

A collection of problems posed or studied by mathematician Paul Erdős, open as of August 2026 and curated to be especially interesting and difficult.

Highest score: 3% | Mathematics | Agent | Epoch administered

View benchmark

MirrorCode

By Epoch

A long-horizon coding benchmark. Models reimplement entire programs end-to-end, without access to the original source code.

Highest score: 73% | Software engineering | Agent | Epoch administered

View benchmark

Earthborne Rangers (EBR-bench)

By Epoch

A test of AI systems' learning capabilities, measuring whether their scores improve across repeated playthroughs of the relatively obscure game Earthborne Rangers.

Highest score: 76% | Games | Agent | Epoch administered

View benchmark

Mystery Game Puzzles

By Epoch

100 positions from a puzzle-oriented variant of a well-known game, the details of which are deliberately left undisclosed.

Highest score: 84% | Games | Epoch administered

View benchmark

Chess Puzzles

By Epoch

100 novel puzzles, generated programmatically with a chess engine. Each puzzle has a single best next move.

Highest score: 72% | Games | Epoch administered

View benchmark

SimpleQA Verified

Verified

1,000 factoid questions about politics, science and technology, art, sports, geography, music, and more.

Highest score: 76% | World knowledge | Epoch administered

Read review | View benchmark

FrontierMath Tier 4 (v2)

By Epoch

Exceptionally difficult research-level math problems.

Highest score: 98% | Mathematics | Agent | Epoch administered

View benchmark

FrontierMath Tiers 1-3 (v2)

By Epoch

Expert-written math problems covering advanced undergrad through early career research.

Highest score: 94% | Mathematics | Agent | Epoch administered

View benchmark

See more

EPOCH AI

Investigating the trajectory of AI for the benefit of society

Get the latest from Epoch AI in your inbox

Subscribe

Our Work
Latest
AI Trends & Statistics
Publications
Data Insights
Newsletter: Gradient Updates
Podcast

Data on AI
Models
AI Data Centers
Hardware
Companies
Chip Sales
Polling

Benchmarking
Capabilities & Benchmarking
FrontierMath

Company
About
Team
Consultations
Transparency
Donate
Careers

© 2026 Epoch AI